

Lecture 11: Graph Search

*Harvard SEAS - Fall 2025**2025-10-07*

1 Announcements

- Salil's OH Thu 1-1:45pm SEC 3.327.
- SRE 3 this Thursday: come prepared!
- Midterm Thu 10/16 (closed book). Practice midterms from 2022 and 2024 posted. This week's section will do midterm review!

Recommended Reading:

- Hesterberg–Vadhan Ch. 12
- CLRS 22.2
- Roughgarden II, 8.1–8.2
- Lewis & Zax 13, 16, 17

Where are we in the course trajectory?

2 Graphs

Motivating Problem: Google Maps.

Given a road network, a starting point, and a destination, what is the shortest way to get from the starting point s to the destination t ?

Q: How to model a road network?

A:

Definition 2.1.

Q: What possibilities doesn't this model capture and how might we augment it?

Unless we state otherwise, assume *graph* means a **simple, unweighted, undirected** graph, and a *digraph* means a **simple, unweighted, directed** graph. Graphs are very useful for modelling many kinds of relationships!

Representing Graphs.

Example 2.2. Consider the graph with vertices $V = [9]$ and edges

$$E = \{(0, 1), (0, 2), (1, 3), (2, 1), (2, 4), (3, 4), (3, 5), (4, 5), (6, 7), (7, 8), (8, 6)\}.$$

- Adjacency list representation: For every vertex v , given

- $\text{Nbr}_{\text{out}}[v] =$

- $\text{deg}_{\text{out}}(v) =$

In above example, $\text{Nbr}_{\text{out}} =$

- Using Word-RAM with word length $w \geq \log n$, so vertex name fits in a single word.

3 Shortest Walks

Abstracting a simplified version of the route-finding problem above, we wish to design an algorithm for the following computational problem:

Input: A digraph $G = (V, E)$ and two vertices $s, t \in V$

Output: A *shortest walk* from s to t in G , if any walk from s to t exists

Computational Problem SHORTEST WALK

Let us define precisely what we mean by a *shortest walk*.

Definition 3.1. Let $G = (V, E)$ be a directed graph, and $s, t \in V$.

- A *walk* w from s to t in G is
- A walk in which all vertices are distinct is also called a *path*.
- The *length* of a walk w is $\text{length}(w) =$
- The *distance* from s to t in G is $\text{dist}_G(s, t) =$
- A *shortest walk* from s to t in G is a walk w from s to t with $\text{length}(w) = \text{dist}_G(s, t)$

Q: What algorithm for ShortestWalk is immediate from the definition?

A:

But when can we stop this algorithm to conclude that there is no walk? The following lemma allows us to stop at walks of length $n - 1$:

Lemma 3.2. *If w is a shortest walk from s to t , then all of the vertices that occur on w are distinct (i.e. w is a path).*

Proof.

□

Because of this lemma, the SHORTEST WALK problem is usually referred to as the SHORTEST PATH problem.

Q: With this lemma, what is the runtime of exhaustive search?

A:

4 Breadth-First Search

“I don’t know where I’m going, but I’m on my way.” — George Fairman

We can get a much faster algorithm for SHORTEST WALK using *breadth-first search* (*BFS*). For simplicity, we’ll start by presenting algorithms to only compute the *length* of the shortest path from s to t , rather than actually find the path. On the other hand, our algorithm will actually compute the distance from s to *all* vertices in the graph, not only t . Let’s capture these two modifications in the following definition:

Input: A digraph $G = (V, E)$ and a source vertex $s \in V$

Output: The array dist_s where for every $t \in V$, $\text{dist}_s[t] = \text{dist}_G(s, t)$

Computational Problem SINGLE-SOURCE DISTANCES

With this, here is our first version of BFS.

BFS(G, s):

Input : A digraph $G = (V, E)$ and a source vertex $s \in V$

Output : The array $\text{dist}_s[\cdot] = \text{dist}_G(s, \cdot)$

0 Initialize $\text{dist}_s[t] = \infty$ for all $t \in V$;

1 $S = F = \{s\}$;

2 $\text{dist}_s[s] = 0$;

3 **foreach** $d = 1, \dots, n - 1$ **do**

4 Let $F =$ _____ ;

5 For every $v \in F$, $\text{dist}_s[v] = d$;

6 $S = S \cup F$;

7 **return** dist_s

Algorithm 4.1: BFS for SINGLE-SOURCE DISTANCES

Let’s illustrate the behavior of BFS to the graph from Example 2.2 with $s = 0, t = 4$:

Q: What is happening at every iteration of the loop?

A:

Q: How do we prove correctness?

A:

Q: What is the runtime of the algorithm, in terms of the number of vertices n and the number of edges m ?

A:

Q: How can we implement BFS faster?

A:

Putting all the above together, we obtain:

Theorem 4.1. *$\text{BFS}(G)$ correctly solves SINGLE-SOURCE DISTANCES and can be implemented in time $O(n + m)$, where n is the number of vertices in G and m is the number of edges in G .*

Implementation details.

- In practice, F is stored as a queue and updates are done one vertex u at a time rather than as a ‘batch’.
- The bitvector S is redundant given that we are maintaining dist_s .

5 Finding the Paths

Q: How to actually find shortest *path* from s to t , not just the distances?

A:

Input: A digraph $G = (V, E)$ and a source vertex $s \in V$

Queries: For any query vertex t , return a shortest path from s to t (if one exists).

Abstract Data Type SINGLE-SOURCE SHORTEST PATHS

Theorem 5.1. SINGLE-SOURCE SHORTEST PATHS *has a data structure that, for digraphs with n vertices and m edges, has preprocessing time $O(n + m)$ and the time to answer a query t is $O(\text{dist}_G(s, t))$.*

Given this, we deduce an algorithm for the SHORTEST WALK problem we started with, by running the preprocessor and then querying the data structure once with our given destination vertex t :

Corollary 5.2. SHORTEST WALK (*equivalently, SHORTEST PATH*) *can be solved in time $O(n + m)$.*

It is natural and very useful (e.g. in Google maps) to have data structures for variants of SHORTEST PATH on *dynamic graphs*. There is a rich collection of methods for dynamic graph data structures, which are beyond the scope of this course.

Another extension of BFS is to handle *weighted graphs*. This is Dijkstra's algorithm, and is covered in CS 1240.

6 (Optional) Other Forms of Graph Search

Another very useful form of graph search that you may have seen is *depth-first search* (DFS). We won't cover it in this textbook, but DFS and some of its applications are covered in CS1240.

We do, however, briefly mention a randomized form of graph search, namely *random walks*, and use it to solve the *decision* problem of whether there *exists* a walk from s to t on undirected graphs:

Input: A graph $G = (V, E)$ and vertices $s, t \in V$

Output: YES if there is a walk from s to t in G , and NO otherwise

Computational Problem UNDIRECTED S-T CONNECTIVITY

This problem can be solved by Algorithm 6.1, for an appropriate setting of the

walk length ℓ .

RandomWalk (G, s, t, ℓ):	
Input	: A digraph $G = (V, E)$, a vertices $s, t \in V$, and a walk-length ℓ
Output	: YES or NO
0	$v = s$;
1	foreach $i = 1, \dots, \ell$ do
2	if $v = t$ then return YES;
3	$j = \text{random}(\text{deg}_{\text{out}}(v))$;
4	$v = \text{Nbr}_{\text{out}}[v][j]$;
5	return ∞

Algorithm 6.1: RandomWalk()

Q: What is the advantage of this algorithm over BFS?

A:

It can be shown that if G is an *undirected* graph with n vertices and m edges, then for an appropriate choice of $\ell = O(mn)$, with high probability **RandomWalk**(G, s, t, ℓ) will visit all vertices reachable from s . Thus, we obtain:

Theorem 6.1. *UNDIRECTED S-T CONNECTIVITY can be solved by a Monte Carlo randomized algorithm with arbitrarily small error probability in time $O(mn)$ using only $O(1)$ words of memory in addition to the input (which is not modified during the algorithm).*